

# Credit Card Fraud Detection

Under the guidance of Dr.Cathy Durso

Probability and  
Statistics [4441-1]

## Contents

Introduction.....	1
Data Source .....	1
Overview of Dataset.....	1
Summary Statistics.....	4
Research .....	6
Research Questions .....	6
Methods for addressing research questions. ....	6
Importance of the project .....	8
Data Satisfaction .....	9
Method Applied and Interpretation .....	10
XGBoost.....	10
Visualization.....	10
Logistic Regression.....	13
Visualization.....	13
Conclusion .....	15
Citations .....	16

## Introduction

What is credit card fraud? When someone uses your credit card to buy goods & services or access your personal account without consent is called credit card fraud. In the European Union the credit card fraud in 2013 was approximately €1.44 Billion.

Types of credit card fraud: Some common types of credit card frauds are:

- Card-not-present
- Counterfeit credit-card
- Account or application hack

With the advent of new technology, fraudsters find new ways to scam people and so it is important to learn the signs and act quickly to report suspected frauds.

How to stop credit card fraud? There is a saying ‘Set a thief to catch a thief’, meaning that the best way to catch a thief is to with the help of another thief because both think alike. Hence, to tune thinking like a thief we have tried to implement machine learning models to learn to identify patterns and anomalies of fraudulent transactions from a large data set and flag such transactions in the future.

## Data Source

For our project we have chosen an open-source date-set from Kaggle : <https://www.kaggle.com/datasets/mlg-ulb/creditcardfraud>

## Overview of Dataset

The data set contains a total of 31 variables, seen below, and 284,807 row entries. The data has already been PCA transformed (Dimensionality Reduction), however due to confidentiality issue a lot of the variable names have been masked. As seen below are listed the headers of the variables. Variables “Time” through “Amount” are all dependent variables and “Class” variable is the only dependent variable.

The “Class” dependent variable is labeled “0” for non-fraud transactions and “1” for fraudulent transactions. All the dependent variables are numeric, and the structure of the data can be seen below.

```
[1] "Time"    "V1"      "V2"      "V3"      "V4"      "V5"      "V6"      "V7"
[9] "V8"      "V9"      "V10"     "V11"     "V12"     "V13"     "V14"     "V15"
[17] "V16"     "V17"     "V18"     "V19"     "V20"     "V21"     "V22"     "V23"
[25] "V24"     "V25"     "V26"     "V27"     "V28"     "Amount"  "Class"
```

```
'data.frame':  284807 obs. of  31 variables:
 $ Time   : num  0 0 1 1 2 2 4 7 7 9 ...
 $ V1     : num  -1.36 1.192 -1.358 -0.966 -1.158 ...
 $ V2     : num  -0.0728 0.2662 -1.3402 -0.1852 0.8777 ...
 $ V3     : num  2.536 0.166 1.773 1.793 1.549 ...
 $ V4     : num  1.378 0.448 0.38 -0.863 0.403 ...
 $ V5     : num  -0.3383 0.06 -0.5032 -0.0103 -0.4072 ...
 $ V6     : num  0.4624 -0.0824 1.8005 1.2472 0.0959 ...
```

```

$ V7      : num  0.2396 -0.0788 0.7915 0.2376 0.5929 ...
$ V8      : num  0.0987 0.0851 0.2477 0.3774 -0.2705 ...
$ V9      : num  0.364 -0.255 -1.515 -1.387 0.818 ...
$ V10     : num  0.0908 -0.167 0.2076 -0.055 0.7531 ...
$ V11     : num  -0.552 1.613 0.625 -0.226 -0.823 ...
$ V12     : num  -0.6178 1.0652 0.0661 0.1782 0.5382 ...
$ V13     : num  -0.991 0.489 0.717 0.508 1.346 ...
$ V14     : num  -0.311 -0.144 -0.166 -0.288 -1.12 ...
$ V15     : num  1.468 0.636 2.346 -0.631 0.175 ...
$ V16     : num  -0.47 0.464 -2.89 -1.06 -0.451 ...
$ V17     : num  0.208 -0.115 1.11 -0.684 -0.237 ...
$ V18     : num  0.0258 -0.1834 -0.1214 1.9658 -0.0382 ...
$ V19     : num  0.404 -0.146 -2.262 -1.233 0.803 ...
$ V20     : num  0.2514 -0.0691 0.525 -0.208 0.4085 ...
$ V21     : num  -0.01831 -0.22578 0.248 -0.1083 -0.00943 ...
$ V22     : num  0.27784 -0.63867 0.77168 0.00527 0.79828 ...
$ V23     : num  -0.11 0.101 0.909 -0.19 -0.137 ...
$ V24     : num  0.0669 -0.3398 -0.6893 -1.1756 0.1413 ...
$ V25     : num  0.129 0.167 -0.328 0.647 -0.206 ...
$ V26     : num  -0.189 0.126 -0.139 -0.222 0.502 ...
$ V27     : num  0.13356 -0.00898 -0.05535 0.06272 0.21942 ...
$ V28     : num  -0.0211 0.0147 -0.0598 0.0615 0.2152 ...
$ Amount: num  149.62 2.69 378.66 123.5 69.99 ...
$ Class  : int  0 0 0 0 0 0 0 0 0 0 ...

```

To give a better understanding of the data we are working with, we look at the first 6 rows from the data set.

Time	V1	V2	V3	V4	V5	V6
1	0 -1.3598071	-0.07278117	2.5363467	1.3781552	-0.33832077	0.46238778
2	0 1.1918571	0.26615071	0.1664801	0.4481541	0.06001765	-0.08236081
3	1 -1.3583541	-1.34016307	1.7732093	0.3797796	-0.50319813	1.80049938
4	1 -0.9662717	-0.18522601	1.7929933	-0.8632913	-0.01030888	1.24720317
5	2 -1.1582331	0.87773675	1.5487178	0.4030339	-0.40719338	0.09592146
6	2 -0.4259659	0.96052304	1.1411093	-0.1682521	0.42098688	-0.02972755

	V7	V8	V9	V10	V11	V12
1	0.23959855	0.09869790	0.3637870	0.09079417	-0.5515995	-0.61780086
2	-0.07880298	0.08510165	-0.2554251	-0.16697441	1.6127267	1.06523531
3	0.79146096	0.24767579	-1.5146543	0.20764287	0.6245015	0.06608369
4	0.23760894	0.37743587	-1.3870241	-0.05495192	-0.2264873	0.17822823
5	0.59294075	-0.27053268	0.8177393	0.75307443	-0.8228429	0.53819555
6	0.47620095	0.26031433	-0.5686714	-0.37140720	1.3412620	0.35989384


	V13	V14	V15	V16	V17	V18
1	-0.9913898	-0.3111694	1.4681770	-0.4704005	0.20797124	0.02579058
2	0.4890950	-0.1437723	0.6355581	0.4639170	-0.11480466	-0.18336127
3	0.7172927	-0.1659459	2.3458649	-2.8900832	1.10996938	-0.12135931
4	0.5077569	-0.2879237	-0.6314181	-1.0596472	-0.68409279	1.96577500
5	1.3458516	-1.1196698	0.1751211	-0.4514492	-0.23703324	-0.03819479
6	-0.3580907	-0.1371337	0.5176168	0.4017259	-0.05813282	0.06865315

	V19	V20	V21	V22	V23	V24
1	0.40399296	0.25141210	-0.018306778	0.277837576	-0.11047391	0.06692807

2	-0.14578304	-0.06908314	-0.225775248	-0.638671953	0.10128802	-0.33984648
3	-2.26185710	0.52497973	0.247998153	0.771679402	0.90941226	-0.68928096
4	-1.23262197	-0.20803778	-0.108300452	0.005273597	-0.19032052	-1.17557533
5	0.80348692	0.40854236	-0.009430697	0.798278495	-0.13745808	0.14126698
6	-0.03319379	0.08496767	-0.208253515	-0.559824796	-0.02639767	-0.37142658
	V25	V26	V27	V28	Amount	Class
1	0.1285394	-0.1891148	0.133558377	-0.02105305	149.62	0
2	0.1671704	0.1258945	-0.008983099	0.01472417	2.69	0
3	-0.3276418	-0.1390966	-0.055352794	-0.05975184	378.66	0
4	0.6473760	-0.2219288	0.062722849	0.06145763	123.50	0
5	-0.2060096	0.5022922	0.219422230	0.21515315	69.99	0
6	-0.2327938	0.1059148	0.253844225	0.08108026	3.67	0

The data set is an unbalanced data set i.e., we have only 492 fraudulent transactions of the total 284,807 transactions that is less than 0.2% of the data set. We can visualize this from the pie chart given below:



## Summary Statistics

The summary statistics are shown below:

Time	V1	V2	V3
Min. : 0	Min. : -56.40751	Min. : -72.71573	Min. : -48.3256
1st Qu.: 54202	1st Qu.: -0.92037	1st Qu.: -0.59855	1st Qu.: -0.8904
Median : 84692	Median : 0.01811	Median : 0.06549	Median : 0.1799
Mean : 94814	Mean : 0.00000	Mean : 0.00000	Mean : 0.00000
3rd Qu.: 139321	3rd Qu.: 1.31564	3rd Qu.: 0.80372	3rd Qu.: 1.0272
Max. : 172792	Max. : 2.45493	Max. : 22.05773	Max. : 9.3826
V4	V5	V6	V7
Min. : -5.68317	Min. : -113.74331	Min. : -26.1605	Min. : -43.5572
1st Qu.: -0.84864	1st Qu.: -0.69160	1st Qu.: -0.7683	1st Qu.: -0.5541
Median : -0.01985	Median : -0.05434	Median : -0.2742	Median : 0.0401
Mean : 0.00000	Mean : 0.00000	Mean : 0.00000	Mean : 0.00000
3rd Qu.: 0.74334	3rd Qu.: 0.61193	3rd Qu.: 0.3986	3rd Qu.: 0.5704
Max. : 16.87534	Max. : 34.80167	Max. : 73.3016	Max. : 120.5895
V8	V9	V10	V11
Min. : -73.21672	Min. : -13.43407	Min. : -24.58826	Min. : -4.79747
1st Qu.: -0.20863	1st Qu.: -0.64310	1st Qu.: -0.53543	1st Qu.: -0.76249
Median : 0.02236	Median : -0.05143	Median : -0.09292	Median : -0.03276
Mean : 0.00000	Mean : 0.00000	Mean : 0.00000	Mean : 0.00000
3rd Qu.: 0.32735	3rd Qu.: 0.59714	3rd Qu.: 0.45392	3rd Qu.: 0.73959
Max. : 20.00721	Max. : 15.59500	Max. : 23.74514	Max. : 12.01891
V12	V13	V14	V15
Min. : -18.6837	Min. : -5.79188	Min. : -19.2143	Min. : -4.49894
1st Qu.: -0.4056	1st Qu.: -0.64854	1st Qu.: -0.4256	1st Qu.: -0.58288
Median : 0.1400	Median : -0.01357	Median : 0.0506	Median : 0.04807
Mean : 0.0000	Mean : 0.00000	Mean : 0.0000	Mean : 0.00000
3rd Qu.: 0.6182	3rd Qu.: 0.66251	3rd Qu.: 0.4931	3rd Qu.: 0.64882
Max. : 7.8484	Max. : 7.12688	Max. : 10.5268	Max. : 8.87774
V16	V17	V18	
Min. : -14.12985	Min. : -25.16280	Min. : -9.498746	
1st Qu.: -0.46804	1st Qu.: -0.48375	1st Qu.: -0.498850	
Median : 0.06641	Median : -0.06568	Median : -0.003636	
Mean : 0.00000	Mean : 0.00000	Mean : 0.000000	
3rd Qu.: 0.52330	3rd Qu.: 0.39968	3rd Qu.: 0.500807	
Max. : 17.31511	Max. : 9.25353	Max. : 5.041069	
V19	V20	V21	
Min. : -7.213527	Min. : -54.49772	Min. : -34.83038	
1st Qu.: -0.456299	1st Qu.: -0.21172	1st Qu.: -0.22839	
Median : 0.003735	Median : -0.06248	Median : -0.02945	
Mean : 0.000000	Mean : 0.00000	Mean : 0.00000	
3rd Qu.: 0.458949	3rd Qu.: 0.13304	3rd Qu.: 0.18638	
Max. : 5.591971	Max. : 39.42090	Max. : 27.20284	
V22	V23	V24	
Min. : -10.933144	Min. : -44.80774	Min. : -2.83663	
1st Qu.: -0.542350	1st Qu.: -0.16185	1st Qu.: -0.35459	
Median : 0.006782	Median : -0.01119	Median : 0.04098	
Mean : 0.000000	Mean : 0.00000	Mean : 0.00000	
3rd Qu.: 0.528554	3rd Qu.: 0.14764	3rd Qu.: 0.43953	

Max. : 10.503090	Max. : 22.52841	Max. : 4.58455
V25	V26	V27
Min. :-10.29540	Min. :-2.60455	Min. :-22.565679
1st Qu.: -0.31715	1st Qu.: -0.32698	1st Qu.: -0.070840
Median : 0.01659	Median : -0.05214	Median : 0.001342
Mean : 0.00000	Mean : 0.00000	Mean : 0.000000
3rd Qu.: 0.35072	3rd Qu.: 0.24095	3rd Qu.: 0.091045
Max. : 7.51959	Max. : 3.51735	Max. : 31.612198
V28	Amount	Class
Min. :-15.43008	Min. : 0.00	Min. :0.000000
1st Qu.: -0.05296	1st Qu.: 5.60	1st Qu.:0.000000
Median : 0.01124	Median : 22.00	Median :0.000000
Mean : 0.00000	Mean : 88.35	Mean :0.001728
3rd Qu.: 0.07828	3rd Qu.: 77.17	3rd Qu.:0.000000
Max. : 33.84781	Max. :25691.16	Max. :1.000000

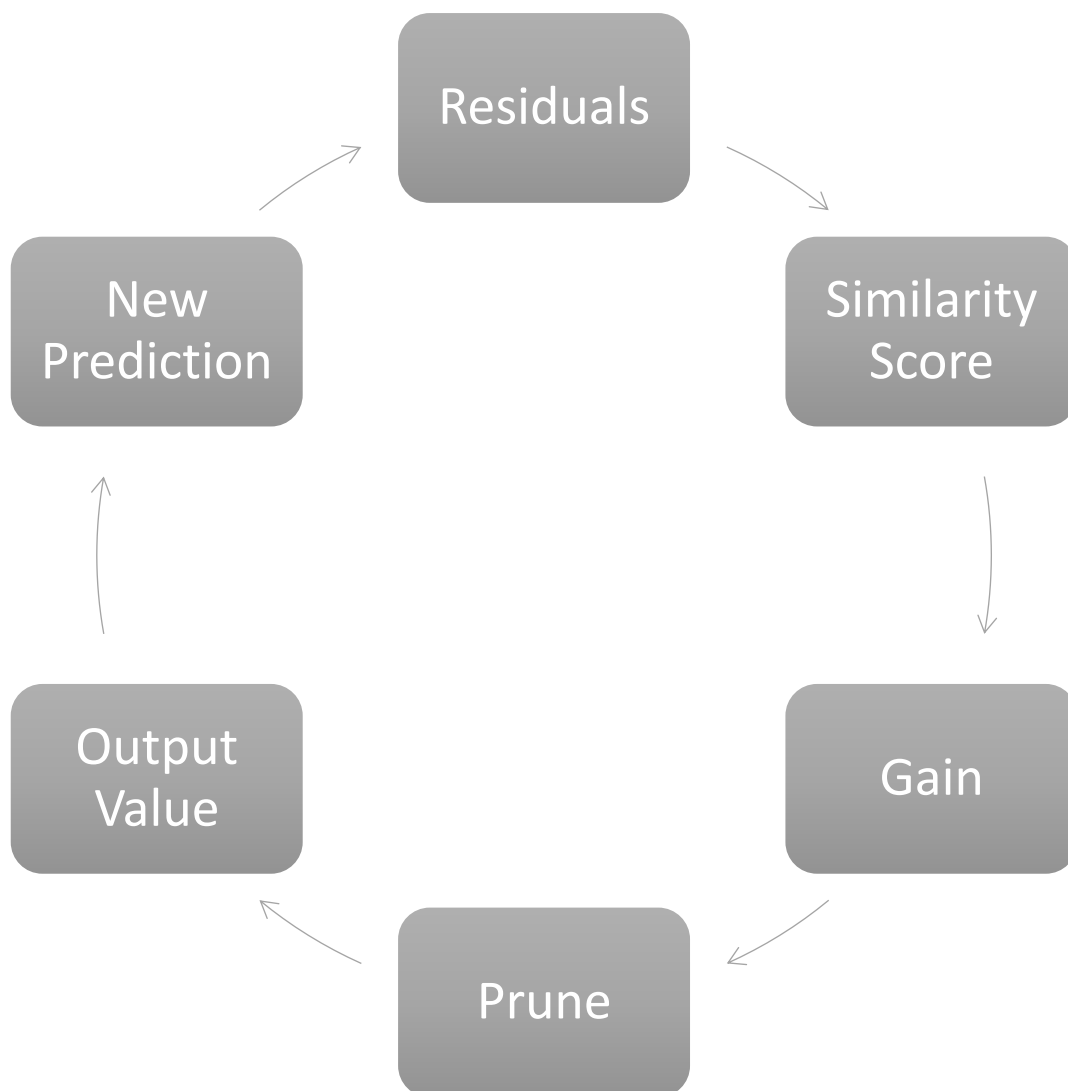
## Research

### Research Questions

- 1) Implement XGBoost and Logistic Regression (for classification) models to predict fraudulent transactions.
- 2) Compare the accuracy of the above models.

### Methods for addressing research questions.

- 1) XGBoost (Classification) - eXtreme Gradient Boost a.k.a XGBoost is a regularized form of gradient boosting. The tool can be used for regression as well as classification. In this project we use the classification method of XGBoost to classify fraudulent transactions from non-fraudulent transactions.





The steps involved in XGBoost are as follows:

- a) Initial prediction - this is usually 0.5 be it for regression or classification.
- b) Similarity Score - this step is a complex step that includes calculating the residuals and then plugging in residual values in the similarity scores formula. We do this for all the leaves combinations i.e. different thresholds. Note: this is iterated until there is only one residual in the Tree or we have achieved tree depth, which is 6 by default.

*SimilarityScore*

$$= \frac{\sum (Residual_i)^2}{\sum [PreviousProbability_i * (1 - PreviousProbability_i)] + \lambda}$$

- c) Gain - To check the clustering of the XGBoost tree, the threshold that gives a higher gain will be used as a branch in the XGBoost tree. Note: this is iterated until there is only one residual in the Tree, or we have achieved tree depth, which is 6 by default.

$$Gain = Left_{similarity} + Right_{similarity} - Root_{similarity}$$

- d) Prune - This is basically cutting the leaves of the tree; the pruning is done based on the gamma value. The gamma value is 0.5 by default. If the difference in gain and gamma is negative, we prune the leaves else leave them as it is.

$$Pruning = Gain - \gamma$$

- e) Output value - After the tree formed, we then calculate the output value with the same lambda as in the similarity score. Note: Lambda is a regularization parameter that reduces the sensitivity of the prediction to isolated observations.

$$Output\ Value = \frac{\sum (Residual_i)}{\sum [PreviousProbability_i * (1 - PreviousProbability_i)] + \lambda}$$

f) New Prediction - Calculated by using the old prediction value, learning rate ( $\epsilon$ , with 0.3 as default value) and the output value. The new prediction residual will be smaller than the residual from the old prediction value.

$$\begin{aligned} \log(\text{odds}) \text{ of New Prediction} \\ = \log(\text{odds})_{\text{old prediction}} + (\epsilon * (\text{Output Value})) \end{aligned}$$

$$\text{New Prediction} = \frac{e^{\log(\text{odds})}}{1 + e^{\log(\text{odds})}}$$

The above steps are iterate until the residuals become very minute or we reach the maximum number of trees.

- 2) Logistic Regression (Classification) - This is like liner regression, but we only use this for classification based on our prediction. The default prediction value is 0.5. We fit the line using maximum likelihood i.e., the line is shifted to evaluate the likelihood and the line with the maximum likelihood is selected.


#### Importance of the project

- 1) Proof of concept for XGBoost classification.
- 2) Model predictor variable as a function of thirty dependent variables to automatically predict fraudulent transactions accurately.
- 3) Comparing accuracy of models in predicting the fraudulent transactions.
- 4) Analyzing significant variables contributing to the predictor variable.

## Data Satisfaction

To carry out analysis there are two important pre-requisites. First, the data must all be numeric. Second, the data needs to be split into Training and Testing sets. The training set will comprise of 80% of the data and will be used to train the machine learning models whereas the testing data will be used to predict the outcome of the “Class” column i.e. 0 for non-fraud transaction and 1 for fraudulent transaction. The testing data is split into two testing sets, one stores the ‘Class’ variable (‘dat.testc’) that will be compared to the other testing data set that will not contain the class variable initially but will be used to predict the transaction in testc under ‘Predicted’ column.

Count of Observations in Training and Testing Data Sets



## Method Applied and Interpretation

### XGBoost

First, we implement the XGBoost (classification) method. XGBoost only accepts matrix as input so we pass the training data set 'dat.train' with the class variable to train the model. The parameters used in the xgboost model are default values such as the eta = 0.3, gamma = 0.5, max\_depth = 6. We did try to tweak the values however we found that the model worked best on these values, giving the maximum accuracy.

### Visualization

#### Confusion Matrix

To compare the outcome

#### Confusion Matrix and Statistics

	Reference	
Prediction	0	1
0	56860	17
1	3	81

Accuracy : 0.9996  
95% CI : (0.9995, 0.9998)  
No Information Rate : 0.9983  
P-Value [Acc > NIR] : < 2e-16

Kappa : 0.8899

Mcnemar's Test P-Value : 0.00365

Sensitivity : 0.9999  
Specificity : 0.8265  
Pos Pred Value : 0.9997  
Neg Pred Value : 0.9643  
Prevalence : 0.9983  
Detection Rate : 0.9982  
Detection Prevalence : 0.9985  
Balanced Accuracy : 0.9132

'Positive' Class : 0

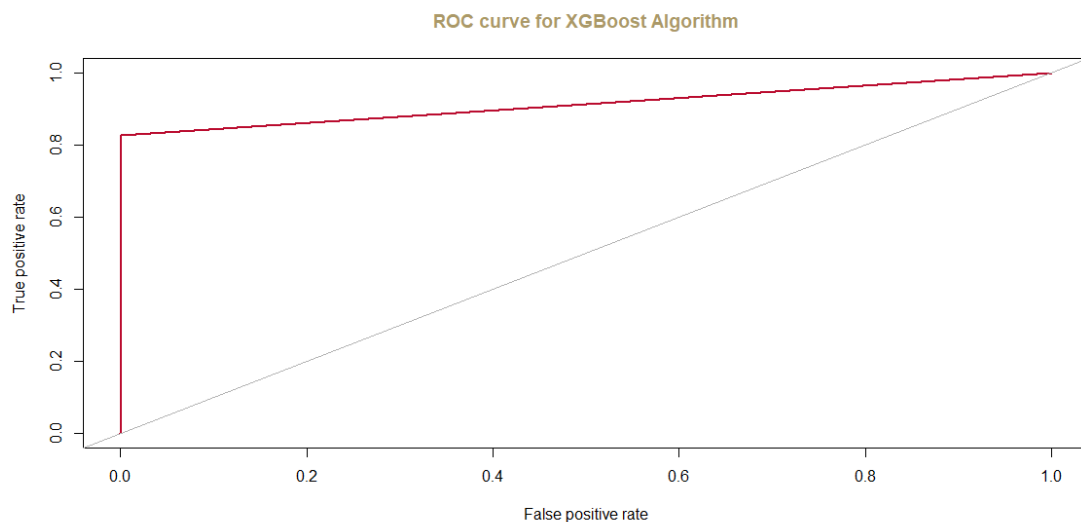
The testing data set contained 56,961 entries. As seen from the confusion matrix above, the XGBoost model for classification correctly identified 99.96% of the transactions, correctly identifying 81 fraudulent transactions and incorrectly marking only 20 transactions. In the incorrect transactions the model incorrectly identified 17 fraudulent transactions as non-fraudulent and 3 non-fraudulent transactions as fraudulent.

Kappa is essentially interrater reliability testing, measure of agreement between the predicted labels and the true labels, and it considers the possibility of agreement occurring by chance. A high Kappa value of

0.8899 means that the classification of this data was not by chance and that the result has almost perfect agreement.

### ROC Curve


The Receiver Operating Characteristic Curve a.k.a ROC Curve, is a graph showing the classification performance of a model at different classification thresholds. The false positive is along the x-axis and the true positives are plotted against the y-axis, and essentially shows the trade-off between clinical sensitivity and specificity. The Area Under the ROC Curve a.k.a AUC provides an cumulative measure of classification performance over possible classification thresholds. The greater the AUC, the higher the ability of the model to distinguish between positive and negative classes.



Area under the curve (AUC): 0.913

The AUC from the XGBoost model is 0.913 which is considered as almost perfect. Moreover, as the goal is to find fraudulent transactions, we can accept a higher false positive rate. Hence, our best threshold will be at the peak of the curve on the top-right corner of the ROC curve.


## Top Contributors



SHAP is an acronym for SHapley Additive exPlanations. SHAP values indicate the contribution of each variable on the final score of the prediction. Seen above are the top 3 variables contributing to the final prediction, the variables are arranged in a descending order. The SHAP values are against the Y-axis and the variable values are against the x-axis. Each blue dot is an entry in the data set, whereas the red curve is the range of values the variable can take and corresponding SHAP values.

Positive SHAP value means positive impact on prediction, leading the model to predict 1.[citation 1]  
Negative SHAP value means negative impact, leading the model to predict 0.[citation 1]

From the graph of variable 'V4' we see that for the range of variable values between 1 through the SHAP values are positive and negative otherwise. This means that variable values of V4 between 1 to 15 have a positive impact, leading the model to predict 1 and predict 0 for other values.



The graph above represents a summary of all the SHAP value of all the 30 independent variables. Each dot on the graph represent an entry in the data set. The heat map on the right-hand side give the range of values that variable takes.

We can see that higher feature value of variable V14 contribute negatively to the prediction. The same can be compared with the 'xgb.plot.shap' and we can see that for V14 for values -1 and greater the SHAP values are negative.

The variables in the graph are in the descending order i.e., the variable V17 contributes the highest in terms of predicting the outcome and V25 contributes the least to the prediction of the outcome.

## Logistic Regression

To compare the results of the above XGBoost Classification model we ran a logistic regression classification to predict the non-fraud and fraudulent cases.

### Visualization

#### Confusion Matrix

##### Confusion Matrix and Statistics

	Reference	
Prediction	0	1
0	56857	29
1	6	69

Accuracy : 0.9994  
95% CI : (0.9991, 0.9996)  
No Information Rate : 0.9983  
P-Value [Acc > NIR] : 1.953e-13

Kappa : 0.7974

Mcnemar's Test P-Value : 0.0002003

Sensitivity : 0.9999  
Specificity : 0.7041  
Pos Pred Value : 0.9995  
Neg Pred Value : 0.9200  
Prevalence : 0.9983  
Detection Rate : 0.9982  
Detection Prevalence : 0.9987  
Balanced Accuracy : 0.8520

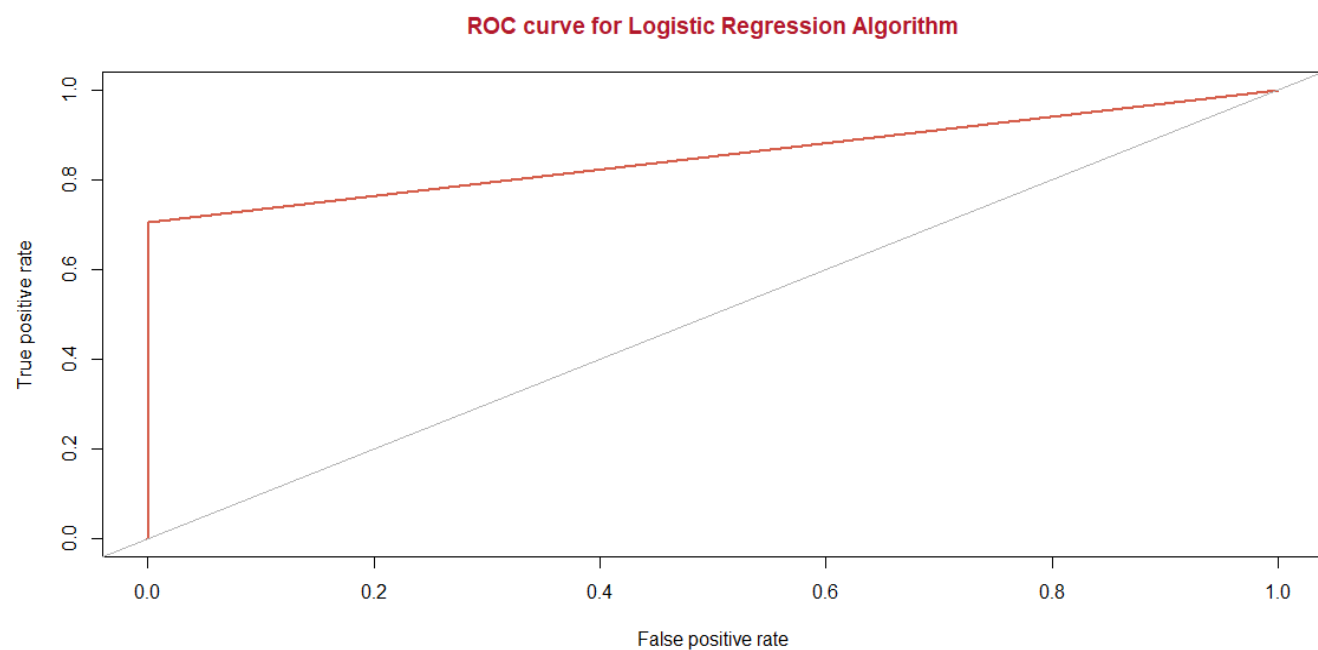
'Positive' Class : 0

The testing data set contained 56,961 entries. As seen from the confusion matrix above, the Logistic Regression model for classification correct identified 99.94% of the transactions, correctly identifying 69 fraudulent transactions and incorrectly marking 35 transactions. In the incorrect transactions the

model incorrectly identified 29 fraudulent transactions as non-fraudulent and 6 non-fraudulent transactions as fraudulent.

Kappa is essentially interrater reliability testing, measure of agreement between the predicted labels and the true labels, and it takes into account the possibility of agreement occurring by chance. A high Kappa value of 0.7974 means that the classification of this data was not by chance and that the result has good agreement.

### ROC Curve



Area under the curve (AUC): 0.852

The AUC from the Logistics Regression (classification) model is 0.852 which is quite high. Moreover, as the goal is to find fraudulent transactions, we can accept a higher false positive rate. Hence, our best threshold will be at the peak of the curve on the top-right corner of the ROC curve.



## Conclusion

Both the XGBoost and logistic regression for classification were implemented on given unbalanced dataset. The findings are:

- 1) The AUC for the XGBoost is significantly better at 0.913 compared to Logistic regression at 0.852, indicating that XGB has better discriminating power.
- 2) The XGBoost model was 42.85% less prone to incorrect classification, which is evident from the confusion matrix of the two model where XGBoost classified 15 fewer transactions incorrectly from a data set of 59,916. Furthermore, we are more concerned about false negatives and on this front the XGBoost model classified 58.6% fewer variables as false negatives.
- 3) The kappa value for XGBoost model and LR model are 0.8899 and 0.7974 respectively, indicating substantial level of agreement between the predicted and true values.

To conclude, the XGBoost (Classification) model for detecting credit card frauds was more robust at correctly predicting fraudulent transactions as compared to Logistic Regression (Classification).

## Citations

1. [Interpretation of SHAP values](#)
2. [Interpretation of SHAP values alternate](#)
3. [XGBoost Mathematics Explained](#)
4. [XGBoost: A Scalable Tree Boosting System](#)
5. [XGBoost Documentation](#)
6. [Interrater reliability: the kappa statistic](#)